


# C++ Threads

*Lawrence Cowl*

*Google*

*13 April 2007*



# introduction

# goals for the standard

- extend the language into concurrency
- interact with the computational environment
- enable expressive libraries for concurrency

# goals for this talk

- outline the primary features
- expose issues under debate
- get feedback

# disclaimers

- talks necessarily suppress details
- every detail in this talk is wrong
- the standard is not done
- its details will change
- compilers will not agree at first

# why add threads?

- concurrency and parallelism is a pressing need
  - effective internet programming
  - effective use of multi-core processors on desktops
  - solution of very large problems
- a standard solution in C++
  - enables portable expression of such programs
  - creates a community of understanding
  - leverages the work of communities

# standardize on the environment

- C++ threads = OS threads
  - heavyweight, pre-emptive, independent
- shared memory
- loosely based on POSIX and Windows
- not a replacement for other standards
  - MPI, OpenMP, automatic parallelization, etc.

# core versus library

- core language changes
  - how do two threads share memory?
  - what operations are atomic?
  - how does this affect variables?
- standard library changes
  - how do programs create and schedule threads?
  - how do threads synchronize and terminate?
  - is that all there is?



memory

# instant shared memory

- all writes are instantly available to all threads
- traditional notion of shared memory
- but there are problems
  - it implies faster-than-light communication
  - it does not match current hardware
  - it inhibits most serial optimizations
- therefore it is not viable

# message shared memory

- writes are explicitly communicated
  - between pairs of threads
  - through a lock or an atomic variable
- the mechanism is acquire and release
  - one thread *releases* its memory writes
    - `v = 32; atomic_store_release( &a, 3 );`
  - another thread *acquires* those writes
    - `i = atomic_load_acquire( &a ); i + v;`

# memory fences

- most shared memory processors have them
- they are not currently in the proposal
- we keep using them in examples
- they imply global action
  - may inhibit more loosely coupled machines
  - may inhibit distributed shared memory
- may or may not find their way into the standard

# sequencing redefined

- sequence points are gone
- sequence now defined by ordering relations
  - sequenced-before
  - indeterminately-sequenced.
- a write/write or read/write pair to a location
  - that are not sequenced-before
  - that are not indeterminately-sequenced
  - results in undefined behavior

# sequencing extended

- sequenced-before
  - provides intra-thread ordering
- acquire and release
  - provide inter-thread ordering
- defining the happens-before relation
  - between memory operations in different threads

# data race condition

- a non-atomic write to a memory location in one thread
- a non-atomic read from or write to that same location in another thread
- with no happens-before relation between them
- is undefined behavior

# memory location

- a non-bitfield primitive data object
- a sequence of adjacent bitfields
  - not separated by a structure boundary
  - not interrupted by the null bitfield
- avoids expensive atomic read-modify-write operations on bitfields



# effect on optimization

- relatively rare optimizations are restricted
  - fewer speculative writes
  - fewer speculative reads
- relatively common optimization have special help
  - they may assume that loops terminate
  - nearly always true

atomics

# atomic definition

- operations
  - on a single variable
  - appear to execute sequentially
- all threads
  - observe the same sequence of values
  - but may not read all those values

# not volatile

- the volatile qualifier has a long history
- we chose not to change its meaning
- volatile remains a "device register" mechanism
- volatile does not indicate atomicity
- an atomic variable can also be volatile

# requirements on atomics

- static initialization
- reasonable implementation on current hardware
- enable (relative) novices to write working code
- enable experts to write very efficient code
- provide a foundation for lock-free data structures

# atomics and memory

- operations specify a memory ordering
  - acquire, release, acq\_rel (both), relaxed (neither)
  - ordered – extra consistency semantics
- too little ordering will break programs
- too much ordering will slow them down
- most programmers should be conservative
  - experts argue about the ordering
  - usually the performance is adequate

# consistency problem

- x and y are atomic and initially 0
  - thread 1:  $x = 1$ ;
  - thread 2:  $y = 1$ ;
  - thread 3: if (  $x == 1 \ \&\& \ y == 0$  )
  - thread 4: if (  $x == 0 \ \&\& \ y == 1$  )
- are both conditions exclusive?
  - that is, is there a total store order?
- the hardware/software system may not provide it
- programming is harder without it

# consistency options

- sequential consistency
  - the observed facts are consistent with a sequential ordering of all events in the system
- weaker models
  - see Hans Boehm's talk, Saturday 9:30
- details and formalism are difficult



# minimal atomics

- atomic flag type
  - static atomic\_flag v1 = ATOMIC\_FLAG\_INIT;
  - if ( atomic\_flag\_test\_set( & v1 ) )
  - atomic\_flag\_clear( & v1 );
- sufficient to implement the rest of the atomics
- similar to patch-panel programming
- users will expect more

# basic atomics

- atomic bool
  - load, store, swap, compare-and-swap
- atomic integers
  - load, store, swap, compare-and-swap,
  - fetch-and-`{ add, sub, and, ior, xor }`
- atomic void pointer
  - load, store, swap, compare-and-swap,
  - fetch-and-`{ add, sub }`

# low-level atomics

- types are opaque
- operations are type-generic function macros in C
- operations overloaded functions in C++
- operations include ordering

```
static atomic_long a = { 1 };  
long t = atomic_load_acquire( & a );  
while ( ! atomic_compare_swap_relaxed( & a, t, t/2 ) );  
atomic_fetch_ior_ordered( & a, 1 );
```

# high-level atomics

- types are classes
- member functions and operators are fully ordered

```
atomic_long a = { 1 };  
long l = a; a = 3;  
while ( ! a.compare_swap( &a, &l, l+1 ) );  
++a; a += 4; a &= 3;
```

# atomic assignment

- default assignment operator is wrong
  - non-atomic load and store
- even atomic load and store is wrong
  - users would expect the whole assignment to be atomic
- correcting the problem is difficult
  - it cannot be disabled in C 90
  - disabling it in C++ 98 breaks C compatibility
  - help (possibly) on the way in C++ 0x
    - N2172 POD's Revisited (Revision 2)
    - N2210 Defaulted and Deleted Functions

# atomic template

- makes an atomic type from a non-atomic type
  - must be bitwise copyable and comparable
- defined specializations for basic types and pointers
- suggested specializations for alignment and size

```
atomic< int * > aip = { 0 };  
aip = ip; aip += 4;  
atomic< gnat > ag = { .... };  
while ( ! ag.compare_swap( &ag, &g, g+4 ) );  
atomic< circus > ac; // works, but not recommended
```

# atomic freedom

- lock-free
  - robust to crashes
  - someone will make progress
- wait-free
  - operations complete in a bounded time
- address-free
  - atomicity does not depend on using the same address

# lock-free atomics

- large atomics have no hardware support
  - necessarily implemented with locks
- locks and signals do not mix
  - must be able to test for lock free
- compile-time macros for basic types
  - always lock-free
  - never lock-free
- run-time function for each type



# wait-free atomics

- hard to implement without direct hardware support
  - resulting programs end up being hardware-specific
  - therefore difficult to write portable programs
- few people seem to care
- property unspecified in standard
  - no requirement for it
  - no query about it

# address-free atomics

- memory may not have a consistent address
  - processes sharing memory may not have the same address for that memory
  - memory may be mapped into the address space twice
- atomic operations must be address-free to work
  - one small tool for inter-process communication
- a lock-free operation is address-free
  - not clear we can say this in a standard way
  - but we will make our intent known

# variables

# affect on variables

- thread-local storage (TLS)
- dynamic initialization of static-duration variables
- destruction of static-duration variables

# adopt thread-local storage

- adopt existing practice
  - 5 vendors with few syntactic variations
- define new storage duration and class
- `__thread int var = 3;`
- variable is unique to each thread
- variable is accessible from every thread
- variable address is not constant

# extend thread-local storage

- existing practice supports only static initialization and trivial destructors
- extend practice to support dynamic initializers and destructors
  - `__thread vector<int> var;`
- carefully define initialization to permit lazy allocation
- operating system support can improve efficiency

# initialization of static-duration variables

- dynamic initialization is tricky
  - no syntax to order most initializations
- without synchronization, potential data races
- with synchronization, potential deadlock
- no consensus yet

# function-local static storage

- initialization implicitly synchronized
  - while not holding any locks
- initialization explicitly synchronized

```
int f() {  
    // synchronized  
    static int u protected = g();  
    static int v = h();  
    // unsynchronized  
    static int w private = i();  
    static nonPOD x private (3);  
}
```



# non-local static storage

- initialization implicitly synchronized
- concurrent initialization enabled
- the initialization may not use a dynamically-initialized object defined outside the translation unit

```
extern vector<int> e;  
vector<int> u; // okay, default initialization  
vector<int> v(u); // okay, within translation unit  
vector<int> w(e); // error, outside translation unit
```

# destruction

- do not destruct
  - programs probably will not be correct in the presence of destruction and threads
  - probably the single largest incompatibility
- do destruct
  - rules to concurrently reverse the concurrent initialization
  - requires cooperative shutdown

launching

# fork and join

- very basic thread class
  - fork a function execution
  - void join operation

```
void f();
```

```
void bar()
{
    std::thread t1(f); // f() executes in separate thread
    t1.join(); // wait for thread t1 to end
}
```

# functors

- may also use function-like objects

```
struct c
{
    void operator()() const;
};

void bar()
{
    std::thread t2((c())); // c() executes in separate thread
    t2.join(); // wait for thread t2 to end
}
```

# scheduling

# scheduling

- limited thread scheduling
  - yield
  - sleep
- standard access to non-standard underlying OS thread handles
  - for detailed control
- query for the hardware concurrency

# synchronization



# mutexes

- exclusive (single reader/writer)
- shared (multiple readers)
- convertible (between exclusive and shared)
- upgradeable (right to become the writer)

# locks

- hold a mutex within a given scope
- represents the mutex acquire/release pair
- release occurs in the destructor for the lock

```
std::upgradable_mutex mut;
```

```
void foo() {  
    std::upgradable_lock< std::upgradable_mutex >  
        read_lock( mut );  
    // ... do read operation  
    if ( /* sometimes need to write */ ) {  
        std::exclusive_lock< std::upgradable_mutex >  
            write_lock( std::move( read_lock ) );  
        // ... do write operation, what was read hasn't changed  
    }  
}
```

# condition variables

- threads may wait on a condition variable
  - giving up their hold on the mutex.
- threads may notify a condition variable
  - notified threads re-acquire the mutex and
  - must reevaluate any condition
- benefits
  - easier to use than events
  - enables the monitor pattern

# buffer example

- conditions represent extreme states

```
class buffer {  
    int head, tail, store[10];  
    std::exclusive_mutex mutex;  
    std::condition_variable not_full, not_empty;  
public:  
    buffer() : head( 0 ), tail( 0 ) { }  
    void insert( int arg ) {  
        std::exclusive_lock< std::exclusive_mutex >  
            scoped( mutex );  
        while ( (head+1)%10 == tail )  
            wake.wait();  
        store[head] = arg;  
        head = (head+1)%10;  
        not_empty.notify();  
    }  
};
```

# termination

# voluntary

- return from outermost function
- some form of cooperative, synchronous, exception-based termination
- strong opposition to asynchronous termination or interrupts
- still debating the details

# exceptions

- when the thread function exits via throw
  - call `std::terminate`?
  - propagate exception to joiner?
  - ignore the exception?
- provide a means to manually propagate
  - `std::exception_ptr saved( std::current_exception() );`
  - `std::exception_ptr copied( std::copy_exception( saved ) );`
  - `rethrow_exception( copied );`

# cancellation

- one thread requests another to cancel itself
- that cancellation will create an exception in the target thread
- when it is ready for one
  - when cancellations have not been disabled
  - only at certain synchronous cancellation points
- or it can test for a pending cancellation
- a thread can ignore the request



# cancellation points

- `void std::this_thread::cancellation_point()`
- `void std::this_thread::sleep( std::nanoseconds )`
- `void std::thread::join()`
- `void std::thread::operator>()()`
- `void std::condition<Lock>::wait( Lock& )`
- `template<class Predicate> void std::condition<Lock>::wait( Lock&, Predicate )`
- `bool std::condition<Lock>::timed_wait( Lock&, std::nanoseconds )`
- `template<class Predicate> void std::condition<Lock>::timed_wait( Lock&, std::nanoseconds, Predicate )`

and beyond

# high-level later

- some work is (probably) being deferred to TR2
  - thread pools, groups, ...
  - value-based joins, futures, ...
  - parallel iterators, ...
- reasons for deferral add up
  - lack of pre-existing implementations
  - lack of solid definitions
  - lack of time to provide them

# success?

- we can build the high-level TR2 facilities in a pure library
- which means you can build even higher-level facilities as well

# futures as an example

- a future executes a function
  - making return value available later
  - propagating exceptions to joiners
- technology is
  - return values via N2096
  - exception propagation via N2179
  - cancellation via handles in N2178

# missing lambda

- anonymous nested functions
  - are not yet in the standard
  - enable user-defined control constructs
  - that put synchronization and concurrency together
  - out of the way of computation
  - and separate the tasks of algorithms specification and performance tuning
- we may get lambda
  - time is running out

# conclusions

# the basics are on track

- memory model
- atomic operations
- non-automatic variables
- thread launching and synchronization



# some features need work

- destruction of static-duration variables
- exception propagation
- cancellation
- high-level facilities
- lambda

# the real value comes later

- the standard will define the basics
- the standard will provide the means to extend the basics
- the users will write some really great high-level facilities

# more information

- C++ standard website
  - <http://www.open-std.org>
  - WG21 - C++
  - WG papers
  - 2007
  - N2169
- questions?

